Modeling the price of a diamond

Kenny Oh

Given the following attributes of a diamond, can we predict its price?

- price price in US dollars ($326--$18,823)
- carat weight of the diamond (0.2--5.01)
- cut quality of the cut (Fair, Good, Very Good, Premium, Ideal)
- color diamond colour, from J (worst) to D (best)
- clarity a measurement of how clear the diamond is (I1 (worst), SI2, SI1, VS2, VS1, VVS2, VVS1, IF (best))
- x length in mm (0--10.74)
- y width in mm (0--58.9)
- z depth in mm (0--31.8)
- depth total depth percentage = z / mean(x, y) = 2 * z / (x + y) (43--79)
- table width of top of diamond relative to widest point (43--95)

The results:

Good (simple linear regression on all features)

```
Training Root Mean Squared Error: 1452.589917565248
Testing Root Mean Squared Error: 1454.9427398770472
```

Better (after generating polynomial features)

```
Training Root Mean Squared Error: 1386.0558094156422
Testing Root Mean Squared Error: 1935.3005119898032
```

Best (after using k best feature selection method and generating dummy variables for categorical data [color, clarity, cut])

```
Training Root Mean Squared Error: 1145.6011338817127
Testing Root Mean Squared Error: 1143.0533608583735
```

Looking at the data...

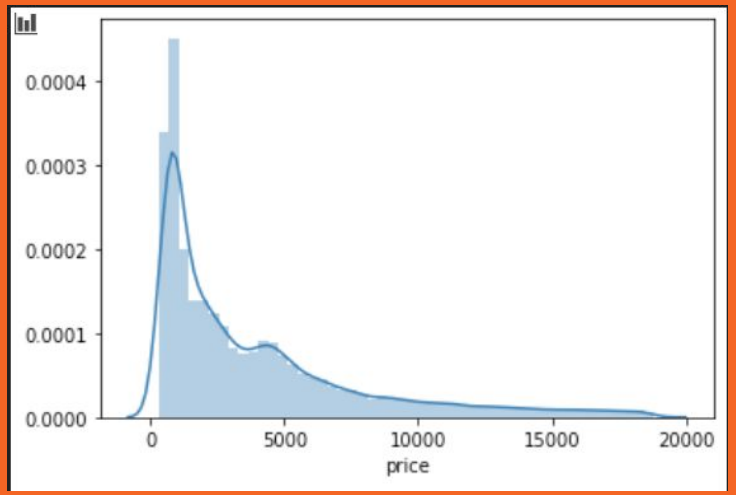

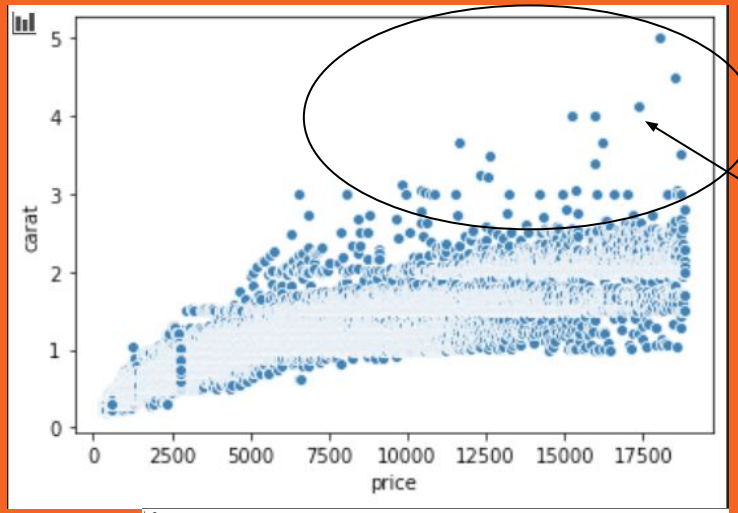

Outliers

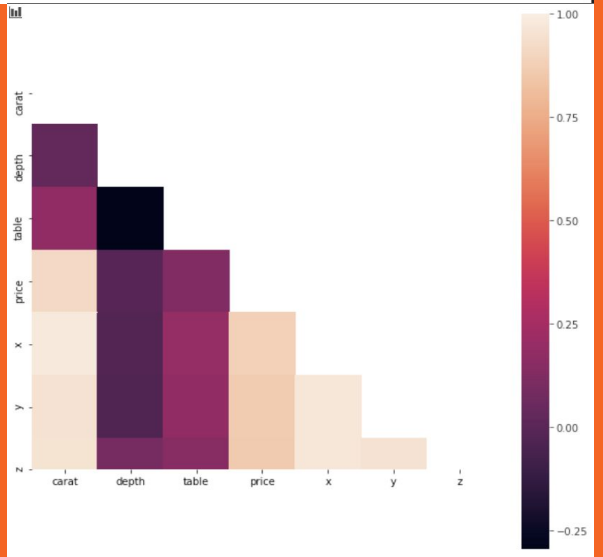

SelectKBest using different scoring functions

f_regression

	Specs	Score
0	carat	313062.636366
11	clarity SI2	876.163036
2	table	872.482672
17	color E	550.912411
5	cut Ideal	502.723390
21	color I	495.598172
6	cut Premium	495.440214
14	clarity VVS1	490.035554
22	color J	342.324716
16	color D	281.369726
20	color H	188.995238
15	clarity VVS2	144.650401
9	clarity IF	131.304090
18	color F	28.805673
3	cut_Fair	13.164617

```
Training Root Mean Squared Error: 1243.5083659785682
Testing Root Mean Squared Error: 1244.4503857632928
```

chi_squared

	Specs	Score
22	color J	18873.657559
11	clarity SI2	18140.426051
9	clarity IF	17589.854955
21	color I	16476.612168
16	color D	16452.629923
15	clarity VVS2	16263.805632
13	clarity VS2	15662.721065
20	color H	15536.255470
10	clarity SI1	15336.163443
14	clarity VVS1	15277.354430
12	clarity VS1	14900.451935
18	color F	14825.423804
17	color E	14600.245627
19	color G	14294.674645
0	carat	14239.086184

```
Training Root Mean Squared Error: 1145.6011338817127
Testing Root Mean Squared Error: 1143.0533608583735
```

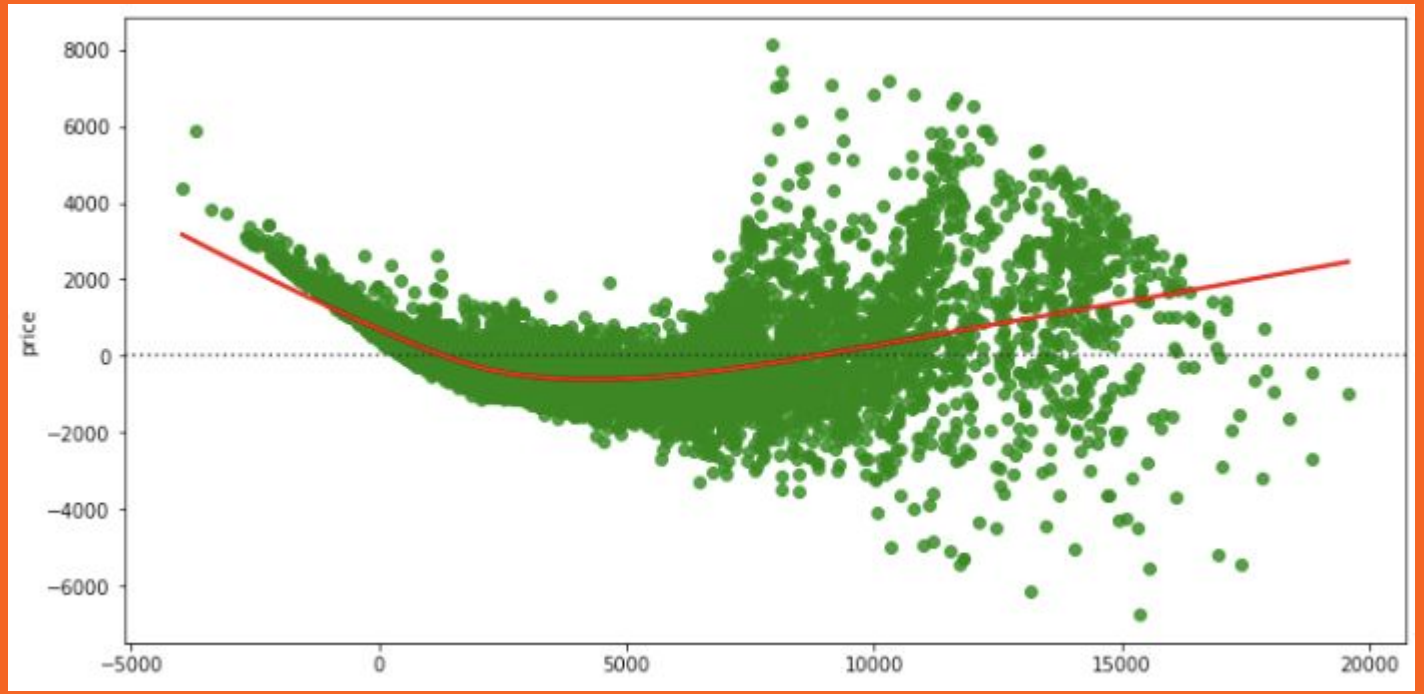
price
8000
6000
4000
2000
0
−2000
−4000
−6000
−5000
0
5000
10000
15000
20000

Thanks!